

Lab: Building a Custom Kubernetes Scheduler in Python

Learning Objectives

By the end of this lab, you will:

- Understand how the **scheduler** interacts with the **API Server**.
 - Implement a **custom scheduler** that:
 - Finds **Pending Pods** with a given **schedulerName**.
 - Chooses a **Node** according to a scheduling policy.
 - **Binds** the Pod to that Node through the Kubernetes API.
 - Compare **polling vs event-driven (watch)** models.
 - Deploy your scheduler into a **kind cluster** and observe its behavior.
-

Environment Setup

Prerequisites

- Docker (latest)
- kind (`go install sigs.k8s.io/kind@latest`)
- kubectl
- Python 3.11+
- (Optional) VS Code or PyCharm for editing

Create and Inspect Cluster

```
kind create cluster --name sched-lab
kubectl cluster-info
kubectl get nodes
```

Step 1 – Observe the Default Scheduler

1. Identify the running scheduler:

```
kubectl -n kube-system get pods -l component=kube-scheduler
kubectl -n kube-system logs -l component=kube-scheduler
```

2. Schedule a simple pod:

```
kubectl run test --image=nginx --restart=Never  
kubectl get pods -o wide
```

Checkpoint 1:

Describe the path:

`kubectl run` → Pod created → Scheduler assigns Node → kubelet starts Pod.

Step 2 — Project Setup

Initialize Project

```
mkdir py-scheduler && cd py-scheduler  
python -m venv .venv && source .venv/bin/activate  
pip install kubernetes==29.0.0  
touch scheduler.py
```

Directory Structure

```
py-scheduler/  
└── scheduler.py  
└── Dockerfile  
└── rbac-deploy.yaml  
└── test-pod.yaml  
└── requirements.txt
```

Step 3 — Implement the Polling Scheduler

Create a simple scheduler in **scheduler.py**:

```
from kubernetes import client, config  
import time, math  
  
def load_client():  
    config.load_incluster_config()  
    return client.CoreV1Api()  
  
def bind_pod(api, pod, node):  
    target = client.V1ObjectReference(kind="Node", name=node)  
    meta = client.V1ObjectMeta(name=pod.metadata.name)  
    body = client.V1Binding(target=target, metadata=meta)  
    api.create_namespaced_binding(pod.metadata.namespace, body)  
  
def choose_node(api):
```

```

nodes = api.list_node().items
pods = api.list_pod_for_all_namespaces().items
min_cnt, pick = math.inf, nodes[0].metadata.name
for n in nodes:
    cnt = sum(1 for p in pods if p.spec.node_name == n.metadata.name)
    if cnt < min_cnt:
        min_cnt, pick = cnt, n.metadata.name
return pick

def main():
    api = load_client()
    while True:
        pods =
api.list_pod_for_all_namespaces(field_selector="spec.nodeName=").items
        for pod in pods:
            if pod.spec.scheduler_name != "my-scheduler":
                continue
            node = choose_node(api)
            bind_pod(api, pod, node)
            print(f"Bound {pod.metadata.name} -> {node}")
            time.sleep(2)

if __name__ == "__main__":
    main()

```

Checkpoint 2:

Understand the control loop:

- **Observe:** list unscheduled Pods
- **Decide:** pick a Node
- **Act:** bind the Pod

Step 4 — Build and Deploy

Dockerfile

```

FROM python:3.11-slim
WORKDIR /app
COPY . .
RUN pip install kubernetes==29.0.0
ENTRYPOINT ["python", "scheduler.py"]

```

Build & Load Image

```

docker build -t my-py-scheduler:latest .
kind load docker-image my-py-scheduler:latest --name sched-lab

```

🔒 Step 5 — RBAC & Deployment

Create a file named **rbac-deploy.yaml**:

```
apiVersion: v1
kind: ServiceAccount
metadata:
  name: my-scheduler
  namespace: kube-system
---
apiVersion: rbac.authorization.k8s.io/v1
kind: ClusterRoleBinding
metadata:
  name: my-scheduler-binding
roleRef:
  apiGroup: rbac.authorization.k8s.io
  kind: ClusterRole
  name: system:kube-scheduler
subjects:
- kind: ServiceAccount
  name: my-scheduler
  namespace: kube-system
---
apiVersion: apps/v1
kind: Deployment
metadata:
  name: my-scheduler
  namespace: kube-system
spec:
  replicas: 1
  selector:
    matchLabels: {app: my-scheduler}
  template:
    metadata:
      labels: {app: my-scheduler}
    spec:
      serviceAccountName: my-scheduler
      containers:
        - name: scheduler
          image: my-py-scheduler:latest
```

Apply it:

```
kubectl apply -f rbac-deploy.yaml
kubectl -n kube-system get pods -l app=my-scheduler
```

✍ Step 6 — Test Your Scheduler

Create a pod definition **test-pod.yaml**:

```
apiVersion: v1
kind: Pod
metadata:
  name: test-pod
spec:
  schedulerName: my-scheduler
  containers:
  - name: pause
    image: registry.k8s.io/pause:3.9
```

Deploy and observe:

```
kubectl apply -f test-pod.yaml
kubectl get pods -o wide
kubectl -n kube-system logs deploy/my-scheduler
```

Checkpoint 3:

Your scheduler should log a message like:

Bound default/test-pod -> kind-control-plane

Step 7 – Event-Driven Scheduler (Watch API)

Replace the polling loop with a **watch-based** approach:

```
import argparse, math
from kubernetes import client, config, watch

# TODO: load_client(kubeconfig) -> CoreV1Api
#   - Use config.load_incluster_config() by default, else
#     config.load_kube_config()

# TODO: bind_pod(api, pod, node_name)
#   - Create a V1Binding with metadata.name=pod.name and
#     target.kind=Node,target.name=node_name
#   - Call api.create_namespaced_binding(namespace, body)

# TODO: choose_node(api, pod) -> str
#   - List nodes and pick one based on a simple policy (fewest running
#     pods)

def main():
    parser = argparse.ArgumentParser()
    parser.add_argument("--scheduler-name", default="my-scheduler")
    parser.add_argument("--kubeconfig", default=None)
    args = parser.parse_args()
```

```

# TODO: api = load_client(args.kubeconfig)

print(f"[watch-student] scheduler starting... name={args.scheduler_name}")
w = watch.Watch()
# Stream Pod events across all namespaces
for evt in w.stream(client.CoreV1Api().list_pod_for_all_namespaces,
_request_timeout=60):
    obj = evt['object']
    if obj is None or not hasattr(obj, 'spec'):
        continue
    # TODO: Only act on Pending pods that target our schedulerName
    # - if obj.spec.node_name is not set and obj.spec.scheduler_name
== args.scheduler_name:
        #     node = choose_node(api, obj)
        #     bind_pod(api, obj, node)
        #     print(...)

if __name__ == "__main__":
    main()

```

Checkpoint 4:

Compare responsiveness and efficiency between **polling** and **watch** approaches.

Step 8 — Policy Extensions

1. Label-based node filtering

```

nodes = [n for n in api.list_node().items
         if "env" in (n.metadata.labels or {}) and
         n.metadata.labels["env"] == "prod"]

```

2. **Taints and tolerations** Use `node.spec.taints` and `pod.spec.tolerations` to filter nodes before scoring.

3. **Backoff / Retry** Use exponential backoff when binding fails due to transient API errors.

4. **Spread policy** Distribute similar Pods evenly across Nodes.

Checkpoint 5:

Demonstrate your extended policy via pod logs and placement.

Observability and Debugging

Command	Description
---------	-------------

Command	Description
<code>kubectl get pods --o wide</code>	Check node assignments
<code>kubectl get events --sort-by=.lastTimestamp</code>	Watch scheduling events
<code>kubectl describe pod test-pod</code>	Inspect scheduling status
<code>kubectl -n kube-system logs deploy/my-scheduler -f</code>	View scheduler logs

🧹 Cleanup

```
kubectl delete -f test-pod.yaml
kubectl delete -f rbac-deploy.yaml
kind delete cluster --name sched-lab
```

RECEIPT Deliverables

Each group (3 students) must submit:

1. **Project files** (`scheduler.py`, `Dockerfile`, `rbac-deploy.yaml`, `test-pod.yaml`). A tgz git repo is preferred.
2. **Logs/screenshots** showing successful scheduling.
3. A short report (~1 page) describing:
 - Node selection logic.
 - Observations of scheduling behavior.
 - Differences between polling and watch models.

📋 Evaluation Criteria

Criterion	Weight	Description
Working scheduler	30%	Scheduler deploys and binds pods correctly
Correct binding logic	25%	Only handles pods with correct schedulerName
Policy extension	25%	Custom filtering/scoring logic implemented
Reflection/report quality	20%	Explains results and design choices clearly

🧠 Reflection Discussion

- Why is it important that your scheduler writes a **Binding** object instead of patching a Pod directly?
- What are the trade-offs between polling vs event-driven models?
- How do **taints and tolerations** interact with your scheduling logic?
- What are real-world policies you could implement using this framework?

Bonus Exercise (Optional)

Implement a **taint-aware scheduler extension** using this helper:

```
def node_tolerates_taints(node, pod):
    taints = node.spec.taints or []
    tolerations = pod.spec.tolerations or []
    if not taints:
        return True
    for taint in taints:
        tolerated = any(
            tol.key == taint.key and
            (tol.effect == taint.effect or tol.effect is None) and
            (tol.operator == "Exists" or tol.value == taint.value)
            for tol in tolerations
        )
        if not tolerated:
            return False
    return True
```

 End of Lab